Toxic Comment Classification

Aruna Neervannan

Abstract

In this project I have built multi-headed model that is capable of detecting different types of of toxicity like threats, obscenity, insults, and identity-based hate better than Perspectives current models. I am using the kaggle dataset which has comments from Wikipedias talk page edits.

KeywordsToxic Comment, Neural Network, CNN, Bi-directional LSTM

1 Introduction

A flurry of content is produced on a daily basis throughout the Internet, on social media sites, blogs, news websites etc from the online interaction between users. While this situation contributes significantly to the quality of human life, unfortunately it involves enormous dangers, since online texts with high toxicity can cause personal attacks, online harassment and bullying behaviors Discussing things you care about can be difficult.

Labeling is a crucial task that has several applications including targeting "trolls" on forums, preventing bullying on social networks or selecting the most accurate comments on a recommendation engine. Automatic abusive language detection is a difficult but important task for online social media. Fighting abusive language online is becoming more and more important in a world where social media plays a significant role in shaping peoples minds (Perse and Lambe, 2016).

Warner and Hirschberg (2012) and Burnap and Williams (2015) are one of the early researches to use machine learning based classifiers for detecting abusive language. Djuric et al., (2015) incorporated representation word embeddings (Mikolov et al., 2013). Nobata et al. (2016) combined predefined language elements and word embedding to train a regression model. Waseem (2016) used logistic regression with n-grams and user-specific features such as gender and location. Davidson et al. (2017) conducted a deeper investigation on different types of abusive language. Badjatiya et al. (2017) experimented with deep learning-based models using ensemble gradient boost classifiers to perform multiclass classification on sexist and racist language. All approaches have been on one step.

2 Purpose

The purpose of this project is to build models that best classifies the comments into the following categories: toxic, severe_toxic, obscene, threat, insult, identify_hate.

3 Motivation

Traditionally Recurrent Neural Networks are best known to work with text data and Convolutional Neural Networks with image processing. Recently, Convolutional Neural Networks (CNN) are being applied to text classification or natural language processing both to distributed as to discrete embedding of words [1], without using syntactic or semantic knowledge of a language [3]. Also, a recurrent CNN model was proposed recently for text classification without human-designed features [2] by succeeding to outperform both the CNN model as well as other well-established classifiers.

4 DATA

4.1 Source

Data for this project are a large number of Wikipedia comments that is part of the Kaggle dataset.

4.2 Variables

The dataset has the commentID, comment and the toxicity. The comments have been labeled by human raters for toxic behavior. The types of toxicity: **toxic, severe_toxic, obscene, threat, insult, identity_hate** The dataset is composed of 300,000 sentences taken from the comments sections of Wikipedia, hlf of which being labeled into 6 categories. The dataset is highly unbalanced, with between 0.2% and 9% of toxic comments depending on the label being considered.

4.3 Test-Train Data Split

Kaggle Dataset did have a label test and train data set.

4.4 Preprocessing

Keras provides the Tokenizer class for preparing text documents for deep learning. The Tokenizer must be constructed and then fit on either raw text documents or integer encoded text documents. The tokenized input sentences are then padded to a max length of 170.

5 Embeddings

A word embedding is an approach to provide a dense vector representation of words that capture something about their meaning. It turns text into numbers and helps building a low-dimensional vector representation from corpus of text, which preserves the contextual similarity of words.

Word embeddings are an improvement over simpler bag-of-word model word encoding schemes like word counts and frequencies that result in large and sparse vectors (mostly 0 values) that describe documents but not the meaning of the words.

Word embeddings work by using an algorithm to train a set of fixed-length dense and continuous-valued vectors based on a large corpus of text. Each word is represented by a point in the embedding space and these points are learned and moved around based on the words that surround the target word.

It is defining a word by the company that it keeps that allows the word embedding to learn something about the meaning of words. The vector space representation of the words provides a projection where words with similar meanings are locally clustered within the space.

The use of word embeddings over other text representations is one of the key methods that has led to breakthrough performance with deep neural networks on problems like machine translation.

For this project Stanford GloVe embeddings were used: Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): glove.840B.300d.zip

6 Model Architecture

6.1 Logistic Regression

Implemented logistic regression using SciPy libraries and calculated the Cross-validation scores for each of the comment classifications.

Formally, the model logistic regression model is that

$$log(\frac{p(x)}{1-p(x)}) = \beta_0 + x.\beta$$

Solving for p, this gives

$$p(x;b,w) = \frac{e^{\beta_0+x.\beta}}{1+e^{\beta_0+x.\beta}} = \frac{1}{1+e^{-(\beta_0+x.\beta)}}$$

To minimize the mis-classification rate, we should predict Y = 1 when p $\geq$ 0.5 and Y = 0 when p < 0.5. This means guessing 1 whenever β_0 + x. β is non-negative, and 0 otherwise. So logistic regression gives us a linear classifier

6.2 Convolutional Neural Networks

A simple CNN architecture may have been used for multi-label classification of fast text embeddings. CNNs can learn patterns in word embeddings, and as per the dataset of the Toxic Classification Challenge, we can use the sub-word information. After tokenizing words, we need an unsupervised learning algorithm for obtaining vector representations for words:

word embeddings. FastText is the library, in which every word is transformed into a vector with a 1 in its corresponding location. For instance, if our word vector is [hi, how, are, you] and the word we are looking at is you, the input vector for you would just be [0, 0, 0, 1]. This works fine in theory, unless the word vocabulary is huge - and in this case, 210,000 words - which means we would end up with word vectors that consist mainly of a bunch of 0s. Data being preprocessed and embeddings being ready, the model is ready to being build. Here is network architecture using built using tensorflow:

- word embedding

- one-dimensional convolutional layer

- one-dimensional Max Pooling

- another convolutional layer

- one-dimensional global Max Pooling

- a drop-out to prevent over-fitting

- a dense layer

- sigmoid activation

Stochastic Gradient Descent algorithm was used as an optimizer with a learning rate of 1e-4.

This section I implemented the CNN model from Convolutional Neural Networks for Sentence Classification [2]. I used three convolutional layers, with filter size width 128, dense vector dimension 300. The filters width is equal to the vector dimension while their height was 3, 4 and 5, for each convolutional layer respectively. After each convolutional layer a max-over-time pooling operation [3] is applied. The output of the pooling layer concatenate to a fully-connected layer, while the softmax function is applied on final layer.
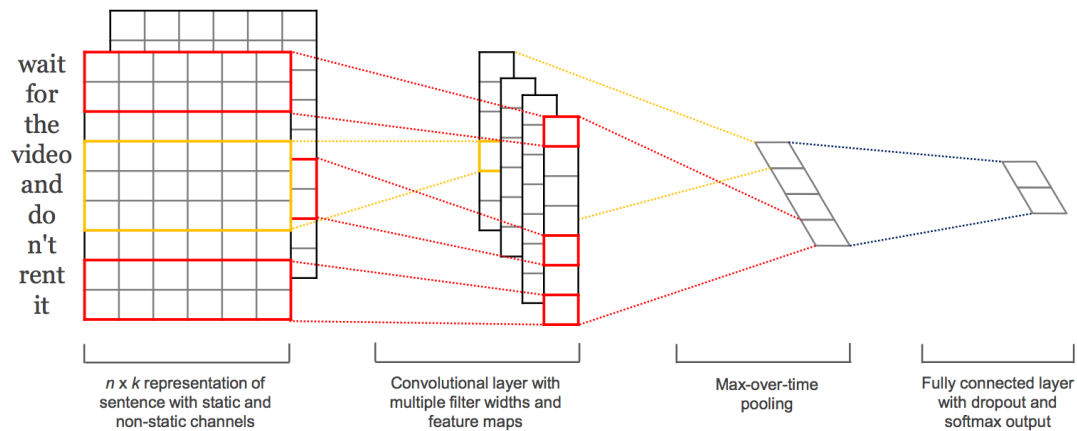

n x k representation of sentence with static and non-static channels	Convolutional layer with multiple filter widths and feature maps	Max-over-time pooling	Fully connected layer with dropout and softmax output

Figure 1: Model architecture with two channels for an example sentence.
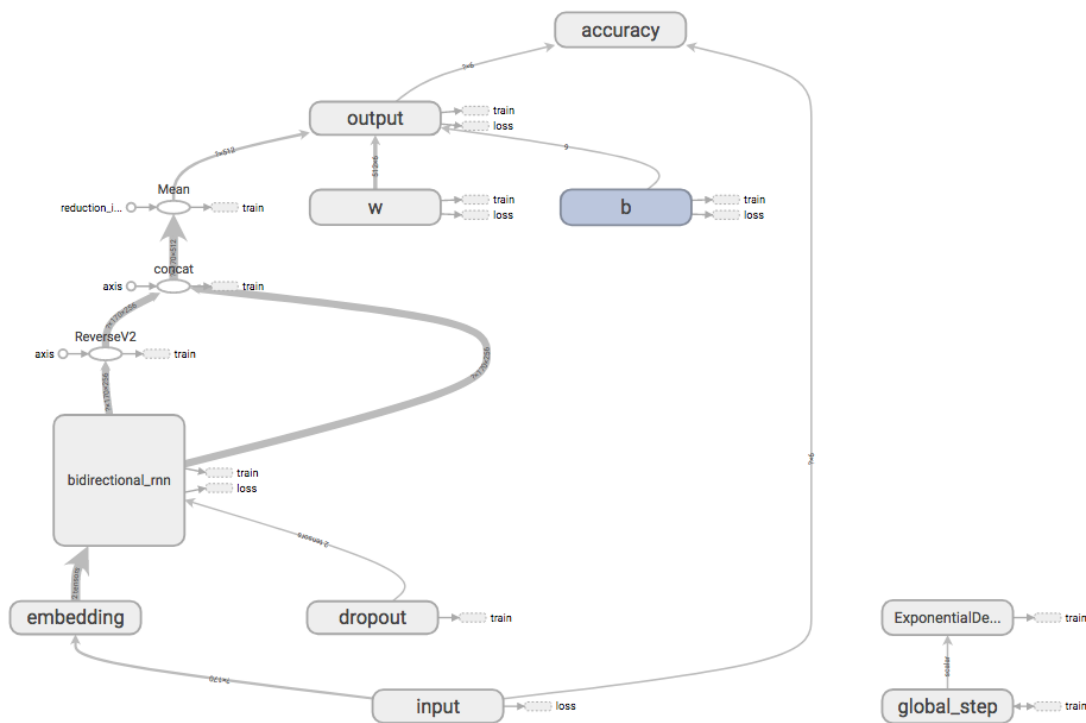
3

Bi-Directional Recurrent Neural Networks

Figure 1: Model architecture for Bi-Directional RNN

Recurrent neural networks (RNN) are networks with loops in them, allowing the persistence of information. A recurrent neural network can be thought of as multiple copies of the same network, each passing a message to a successor. Long Short Term Memory networks, called LSTMs are a special kind of RNN, capable of learning long-term dependencies. All recurrent neural networks have the form of a chain of repeating modules of neural network. In standard RNNs, this repeating module will have a very simple structure, such as a single tanh layer.

One appeal for RNNs is that they might be able to connect previous information to the present task, such as previous sentences might help the understanding of the following sentences. Yet, in reality, it depends. As data length grows, RNNs become un-

able to learn how to connect information. And this is exactly where **Long Short Term Memory neural networks (LSTMs)** come on stage. They are explicitly designed to avoid the long-term dependency problem. The code idea behind LSTMs is the cell state, which is more-or-less like a conveyor belt. It goes from end to end of the entire chain, with only some minor linear interactions. Hence, it is very easy for information to just flow along it without being changed.

The principle of **BiDirectional Recurrent Neural Network (BRNN)** is to split the neurons of a regular RNN into two directions, one for positive time direction (forward states), and another for negative time direction (backward states). Those two states output are not connected to inputs of the opposite direction

states.

Here I have used a stacked Bidirectional LSTM with max k-pooling (a **fully connected feed forward network**). The network here has 500 hidden units in the forward and backward LSTM layers. The output of the pooling layer concatenate to a fully-connected layer, while the softmax function is applied on final layer (which would force the 6 probabilities to add up to 1) with 6 units, which correspond to the predicted probabilities of each of the 6 labels.

Results

6.3 Logistic Regression

Cross Validation scores from Logistic Regression used a baseline scores:

```
CV score for class toxic is -0.141080425752
CV score for class severe_toxic is -0.0297990227241
CV score for class obscene is -0.0804189464205
CV score for class threat is -0.0151034029795
CV score for class insult is -0.0895439792452
CV score for class identity_hate is -0.0330330762226
Total CV score is -0.0648298088908
```

6.4 CNN

Results from training the CNN for 25 epochs:

```
Epoch:0 Train Loss: 0.09688197537347747, Train Accuracy: 0.9730627584610475
Epoch:1 Train Loss: 0.04754850838100164, Train Accuracy: 0.9829976767253723
Epoch:2 Train Loss: 0.03855618966644687, Train Accuracy: 0.9851702522504387
Epoch:3 Train Loss: 0.033110798796994634, Train Accuracy: 0.9871704034016947
Epoch:4 Train Loss: 0.028340381063389358, Train Accuracy: 0.9890294774386129
Epoch:5 Train Loss: 0.024460958881994312, Train Accuracy: 0.9905478791670087
Epoch:6 Train Loss: 0.021070652476917015, Train Accuracy: 0.9917235158228377
Epoch:7 Train Loss: 0.017409704199560096, Train Accuracy: 0.9932398256482512
Epoch:8 Train Loss: 0.01427258154710704, Train Accuracy: 0.9945575834086007
Epoch:9 Train Loss: 0.011801619254976487, Train Accuracy: 0.9955179481789541
Epoch:10 Train Loss: 0.01011599421370805, Train Accuracy: 0.996273490628691
Epoch:11 Train Loss: 0.0083960513797339, Train Accuracy: 0.9969997744690358
Epoch:12 Train Loss: 0.007301356534590948, Train Accuracy: 0.997391653003509
Epoch:13 Train Loss: 0.006378908369844902, Train Accuracy: 0.997729192002053
Epoch:14 Train Loss: 0.005798593860440611, Train Accuracy: 0.9979789488579641
Epoch:15 Train Loss: 0.005340096975951777, Train Accuracy: 0.9982130310317295
Epoch:16 Train Loss: 0.004625524440042175, Train Accuracy: 0.9984241220771214
Epoch:17 Train Loss: 0.00429624308009608, Train Accuracy: 0.998547434041435
Epoch:18 Train Loss: 0.0038189950454832512, Train Accuracy: 0.9987282212434956
Epoch:19 Train Loss: 0.003519605719197555, Train Accuracy: 0.9988295854955768
Epoch:20 Train Loss: 0.003352233339784744, Train Accuracy: 0.9989319963210276
Epoch:21 Train Loss: 0.003245662289891537, Train Accuracy: 0.9989090067234315
```

```
Epoch:22 Train Loss: 0.0029979815607246558, Train Accuracy: 0.9989737982160589
Epoch:23 Train Loss: 0.0029260244573758734, Train Accuracy: 0.9990312729754379
Epoch:24 Train Loss: 0.002874837110806529, Train Accuracy: 0.999025002415174
Training complete
Predictions complete
```

6.5 BRNN

Results from training the BRNN for 25,000 steps:

```
step:24000 Train Loss: 1.3490717, Train Accuracy: 0.9575
step:24100 Train Loss: 1.1947979, Train Accuracy: 0.9609375
step:24200 Train Loss: 1.4230306, Train Accuracy: 0.984375
step:24300 Train Loss: 1.5326409, Train Accuracy: 0.973125
step:24400 Train Loss: 1.3895309, Train Accuracy: 0.9765625
step:24500 Train Loss: 1.2192655, Train Accuracy: 0.96875
step:24600 Train Loss: 1.3316984, Train Accuracy: 0.9765625
step:24700 Train Loss: 1.4945354, Train Accuracy: 0.9775
step:24800 Train Loss: 1.2704673, Train Accuracy: 0.9765625
step:24900 Train Loss: 1.2523654, Train Accuracy: 0.97875
Training complete
Predictions complete
```

7 Summary and Future work

As shown CNN can outperform well established methodologies providing enough evidence that their use is appropriate for toxic comment classification. The promising results are motivating for further development of CNN based methodologies for text mining in the near future, in our interest, employing methods for adaptive learning and providing further comparisons with n-gram based approaches.

While a complex BRNN model produced very similar results it can be clearly seen that CNN can be used to produce similar results without the additional disadvantages with RNN's for eg. increased training time.

Additionally the accuracy can be further increased by techniques such as data augmentation and additionally cleaning the input data and normalizing the text. The embedding and Bi-directional LSTM recurrent states have proven to be great source of information because the symantic information for the toxic phrases/words are well captured and this can be further snapshotted by using multiple layers of ReLU to capture the non-linear behavior on top of such features. Future work in this area would involve hyperparameter tuning on regularizers, adjust the dense dimensions of ReLU layers and calibrating the length of the Bi-directional LSTM layers.

References

[1] Cicero dos Santos and Maira Gatti. 2014. Deep convolutional neural networks for sentiment analysis of short texts. In Proceedings of COLING 2014, the 25th International Conf

[2] Yoon Kim. 2014. Convolutional neural networks for sentence classification. arXiv preprint arXiv:1408.5882 (2014).

[3] Mikael Henaff, Joan Bruna, and Yann LeCun. 2015. Deep Convolutional Networks on Graph-Structured Data. CoRR abs/1506.05163 (2015). arXiv:1506.05163 http://arxiv.org/abs/1506.05163

[4] Badjatiya, P., Gupta, S., Gupta, M., Varma, V.(2017). Deep learning for hate speech detection in tweets. Proceedings of the 26th International Conference on World Wide Web Companion, 759-760.

[5] Djuric, N., Zhou, J., Morris, R., Grbovic, M., Radosavljevic, V., Bhamidipati, N. (2015). Hate speech detection with comment embeddings. Proceedings of the 24th International Conference on World Wide Web, 29-30.

[6] Kingma, D., Ba, J. (2014). Adam: A method for stochastic optimization. Proceedings of the 3rd International Conference on Learning Representations (ICLR)

[7] Wang, W., Chen, L., Thirunarayan, K., Sheth, A.P. (2012). Harnessing twitter" big data" for automatic emotion identification. Privacy, Security, Risk and Trust (PASSAT), 2012 International Conference on and 2012 International Confernece on Social Computing (SocialCom), 587-592.